

# Real-Time Face Re-Identification using YOLO and FaceNet

Shashank Reddy and Prisha Khairnar

## Abstract

This project presents a robust real-time system for face detection, identification, and re-identification in live video streams. The architecture integrates **YOLOv3** for high-speed and accurate face detection, and **FaceNet** for generating compact and discriminative facial embeddings. A tracking mechanism using OpenCV's **CSRT tracker** ensures continuous face tracking across frames. When tracking confidence drops, **re-detection and re-identification** are triggered using cosine similarity and contextual heuristics. The system has been tested under various lighting and occlusion scenarios, demonstrating reliable performance.

## 1. Introduction

Face recognition technologies are increasingly relevant in areas like surveillance, attendance systems, and human-computer interaction. A particular challenge arises in real-time environments: maintaining the identity of a person even after temporary occlusion or disappearance from the camera frame. This project was designed to address that challenge by building an end-to-end pipeline for real-time **face re-identification**.

We started by experimenting with traditional techniques and gradually transitioned to deep learning-based methods to improve robustness and accuracy.

## 2. Initial Approaches

### 2.1 Traditional Methods

We began with conventional approaches that required low computational resources but lacked robustness:

- **Haar Cascades**: Basic detection model; failed for side profiles and non-frontal faces.

- **CLAHE and Histogram Equalization:** Applied for low-light enhancement, but results were inconsistent.
- **Dlib HOG + SVM Detector:** Showed improvement over Haar cascades, but performed poorly with complex lighting and occlusion.

## 2.2 Mediapipe

We experimented with **Google’s Mediapipe face detection** framework, which showed promise due to its lightweight architecture. However, it suffered from significant inconsistencies under varied lighting and head pose conditions, prompting the shift towards deep learning-based detection.

## 3. Final Pipeline

### 3.1 Detection with YOLOv3

We use **YOLOv3**, pretrained on the **WIDER Face dataset**, for real-time detection of human faces. YOLOv3 was selected for its:

- High detection speed (real-time capable)
- Robust performance across a range of facial orientations and lighting conditions

Post-processing includes **Non-Maximum Suppression (NMS)** to eliminate overlapping detections.

### 3.2 Embedding with FaceNet

Initially, we used **Dlib’s facial embedding model**, but it provided subpar accuracy in identity matching. This led us to adopt **FaceNet**, which generates **128-dimensional embeddings** that are both compact and discriminative.

To create a reference for each detected person, embeddings are **averaged over a few frames** and stored for future matching.

### 3.3 Face Matching

Face re-identification is carried out using **cosine similarity** between current and stored embeddings. To enhance robustness, we introduce a **composite similarity score** that combines:

- Cosine similarity (weight: 0.8)
- Positional similarity based on the last known location of the face (weight: 0.2)

### 3.4 Light Adjustment

Handling varying lighting is crucial in real-world environments. Our system includes:

- **Gamma Correction:** Applied dynamically with a gamma value of 1.7
- **Conditional Denoising:** Utilizes **Non-Local Means Denoising** for frames where brightness falls below a predefined threshold

### 3.5 Tracking Mechanism

To maintain the identity of a face across frames, we implemented a tracking system using OpenCV’s **CSRT tracker**. CSRT is robust to scale changes and occlusions, making it suitable for real-time tracking.

Prior to finalizing CSRT, we also experimented with the **Multi-Channel Perception (MCP)** tracker, which was less consistent in handling jitter and motion blur.

The tracking pipeline works as follows:

- A face is selected and tracked using CSRT.
- If the confidence score drops or the tracker fails, the system initiates **re-detection and re-identification**.
- The new candidate face is matched against existing identities using the composite score.

## 4. Re-Detection Heuristics

When a face is lost (e.g., due to occlusion or leaving the frame), re-detection is triggered. To decide whether a newly detected face corresponds to a previously tracked identity, we apply the following heuristics:

- **Brightness Analysis:** The detected face’s brightness is used to set dynamic thresholds for similarity.
- **Composite Score Calculation:** The cosine and positional scores are combined to match the new face to known identities.

- **Temporal Smoothing:** Averaging embeddings over time helps reduce noise and misidentification during re-detection.

## 5. Results

Our system successfully demonstrates real-time performance under varied conditions. Key outcomes include:

- Accurate face detection and tracking across lighting conditions
- Successful re-identification of faces even after brief occlusions
- High reliability in low-light scenarios, due to gamma correction and denoising
- Embedding stability using FaceNet, significantly improving matching over Dlib

## 6. Known Limitations

- If a different person suddenly appears very close to the camera or in front of the tracked face, the tracker may incorrectly switch to the new face.
- Multiple similar-looking faces in a frame can still cause minor misidentifications due to visual overlap.
- The system currently does not use facial landmarks or head pose estimation, which could further improve re-identification accuracy.

## 7. Conclusion

This project successfully implements a real-time face re-identification system using a combination of YOLOv3, FaceNet, and CSRT tracking. The pipeline demonstrates reliability, efficiency, and robustness under varied real-world conditions. Future improvements may include:

- Occlusion-aware models
- Incorporating facial landmarks
- Enhanced tracker switching logic using identity confidence history

## Credits

**Developed by:** Shashank Reddy and Prisha Khairnar

**Motivation:** Final task assigned by Rugved